

	path	label
0	imagesv/v/o/h/voh71d00/509132755+-2755.tif	3
1	imagesl/l/x/t/lxt19d00/502213303.tif	3
2	imagesx/x/e/d/xed05a00/2075325674.tif	2
3	imageso/o/j/b/ojb60d00/517511301+-1301.tif	3
4	imagesq/q/z/k/qzk17e00/2031320195.tif	7

```
In [6]: dir_path = 'data_final'
```

```
In [7]: from sklearn.model_selection import train_test_split
```

```
In [8]: x_train,x_valid=train_test_split(df,test_size=0.30, random_state=42)
```

```
In [9]: datagen=tf.keras.preprocessing.image.ImageDataGenerator(rescale=1./255)
```

```
In [10]: # https://vijayabhaskar96.medium.com/tutorial-on-keras-flow-from-dataframe-1fd4493d237c
train_generator=datagen.flow_from_dataframe(
    dataframe=x_train,
    directory="/content/data_final",
    x_col="path",
    y_col="label",
    batch_size=50,
    seed=42,
    shuffle=True,
    class_mode="categorical",
    target_size=(224,224))
```

Found 33600 validated image filenames belonging to 16 classes.

```
In [11]: # https://vijayabhaskar96.medium.com/tutorial-on-keras-flow-from-dataframe-1fd4493d237c
valid_generator=datagen.flow_from_dataframe(
    dataframe=x_valid,
    directory="/content/data_final",
    x_col="path",
    y_col="label",
    batch_size=50,
```

```
seed=42,  
shuffle=True,  
class_mode="categorical",  
target_size=(224,224))
```

Found 14400 validated image filenames belonging to 16 classes.

Model 1

```
In [12]: from keras.models import Sequential  
from tensorflow.keras.models import Model  
from keras.layers import Dense, Activation, Flatten, Dropout, Input  
from keras.layers import Conv2D, MaxPooling2D  
from keras import regularizers, optimizers  
from keras.applications.vgg16 import VGG16  
from tensorflow.keras.callbacks import ReduceLRonPlateau, EarlyStopping, ModelCheckpoint  
import datetime
```

```
In [ ]: #Use VGG-16 pretrained network without Fully Connected layers and initilize all the weights with Imagenet trained weights  
vgg_ = VGG16(weights='imagenet', include_top=False, input_shape=(224, 224, 3))
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5
58892288/58889256 [=====] - 1s 0us/step

```
In [ ]: vgg_.trainable=False
```

```
In [ ]: vgg_.summary()
```

Model: "vgg16"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0

block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 0		
Non-trainable params: 14,714,688		

```
In [ ]: #After VGG-16 network without FC layers, add a new Conv block ( 1 Conv layer and 1 Maxpooling ), 2 FC layers and a ou
# You are free to choose any hyperparameters/parameters of conv block, FC layers, output layer

#Input layer
input_layer = Input(shape=(224,224,3),name='Input_Layer')
vgg16_layer= vgg_(input_layer)

#vgg16_layer1 =vgg16_layer.layers[-1].output
```

```

#Conv Layer
Conv1 = Conv2D(filters=256,kernel_size=(3,3),strides=(1,1),padding='valid',
               activation='relu',kernel_initializer=tf.keras.initializers.he_normal(seed=23),name='Conv1')(vgg16_layer

#MaxPool Layer
Pool1 = MaxPool2D(pool_size=(2,2),strides=(2,2),padding='valid',name='Pool1')(Conv1)

#Flatten
flatten = Flatten(name='Flatten')(Pool1)
drop=Dropout(0.5)(flatten)

#FC layer
FC1 = Dense(units=512,activation='relu',kernel_initializer=tf.keras.initializers.he_normal(seed=25),name='FC1')(drop)

#FC layer
FC2 = Dense(units=32,activation='relu',kernel_initializer=tf.keras.initializers.he_normal(seed=18),name='FC2')(FC1)

#output layer
Out = Dense(units=16,activation='softmax',kernel_initializer=tf.keras.initializers.glorot_normal(seed=29),name='Output')(FC2)

#Creating a model
model = Model(inputs=input_layer,outputs=Out)
#model1 = Model(inputs = vgg16_layer, outputs = Out)

```

```

In [ ]: #compiling
model.compile(optimizer=tf.keras.optimizers.Adam(lr=0.0001),loss='categorical_crossentropy',metrics=['accuracy'])

```

```

In [ ]: model.summary()

```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
Input_Layer (InputLayer)	[(None, 224, 224, 3)]	0
vgg16 (Functional)	(None, 7, 7, 512)	14714688
Conv1 (Conv2D)	(None, 5, 5, 256)	1179904
Pool1 (MaxPooling2D)	(None, 2, 2, 256)	0
Flatten (Flatten)	(None, 1024)	0

dropout (Dropout)	(None, 1024)	0
FC1 (Dense)	(None, 512)	524800
FC2 (Dense)	(None, 32)	16416
Output (Dense)	(None, 16)	528
=====		
Total params: 16,436,336		
Trainable params: 1,721,648		
Non-trainable params: 14,714,688		

```
In [ ]: %load_ext tensorboard
```

```
In [ ]: !rm -rf ./logs/
import datetime
```

```
In [ ]: import os
logdir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
tensorboard_callback = tf.keras.callbacks.TensorBoard(logdir, histogram_freq=1, write_graph=True, write_grads=True)

WARNING:tensorflow:`write_grads` will be ignored in TensorFlow 2.0 for the `TensorBoard` Callback.
```

```
In [ ]: logdir
```

```
Out[ ]: 'logs/20210510-181438'
```

```
In [ ]: STEP_SIZE_TRAIN=train_generator.n//train_generator.batch_size
STEP_SIZE_VALID=valid_generator.n//valid_generator.batch_size

model.fit(train_generator,
          steps_per_epoch=STEP_SIZE_TRAIN,
          validation_data=valid_generator,
          validation_steps=STEP_SIZE_VALID,
          epochs=6,
          callbacks=[tensorboard_callback])
```

```
Epoch 1/6
672/672 [=====] - 314s 407ms/step - loss: 2.3492 - accuracy: 0.2587 - val_loss: 1.4531 - val
_accuracy: 0.5648
```

```

Epoch 2/6
672/672 [=====] - 265s 394ms/step - loss: 1.4684 - accuracy: 0.5512 - val_loss: 1.2340 - val
_accuracy: 0.6253
Epoch 3/6
672/672 [=====] - 255s 379ms/step - loss: 1.2562 - accuracy: 0.6178 - val_loss: 1.1230 - val
_accuracy: 0.6621
Epoch 4/6
672/672 [=====] - 238s 353ms/step - loss: 1.1442 - accuracy: 0.6536 - val_loss: 1.0753 - val
_accuracy: 0.6747
Epoch 5/6
672/672 [=====] - 228s 339ms/step - loss: 1.0619 - accuracy: 0.6795 - val_loss: 1.0315 - val
_accuracy: 0.6878
Epoch 6/6
672/672 [=====] - 223s 332ms/step - loss: 0.9936 - accuracy: 0.6944 - val_loss: 0.9940 - val
_accuracy: 0.7017

```

Out[]: <tensorflow.python.keras.callbacks.History at 0x7f09d4223a50>

```

In [ ]: %load_ext tensorboard
        %tensorboard --logdir $logdir

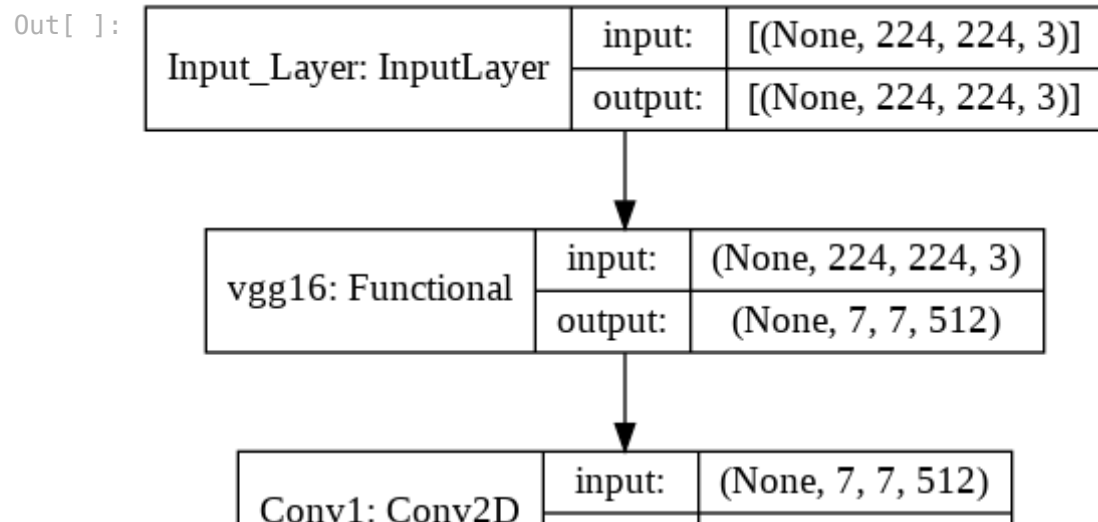
```

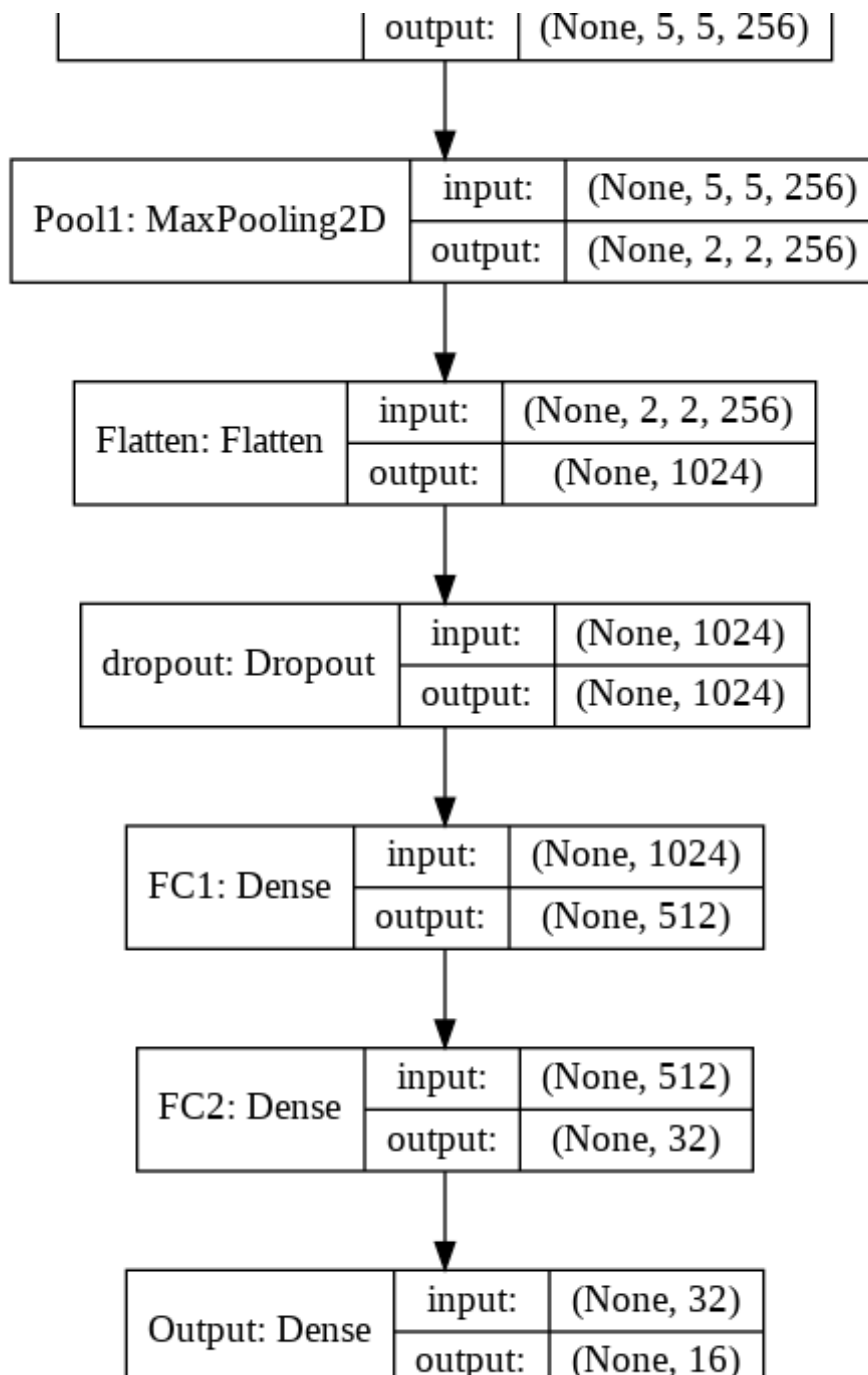
The tensorboard extension is already loaded. To reload it, use:
%reload_ext tensorboard

```

In [ ]: from tensorflow.keras.utils import plot_model
        plot_model(model, 'model.png', show_shapes=True)

```





Model 2

```
In [ ]: tf.keras.backend.clear_session()
```

```
In [ ]: import os
logdir2 = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
tensorboard_callback2 = tf.keras.callbacks.TensorBoard(logdir2, histogram_freq=1, write_graph=True, write_grads=True)
```

WARNING:tensorflow:`write_grads` will be ignored in TensorFlow 2.0 for the `TensorBoard` Callback.

```
In [ ]: logdir2
```

```
Out[ ]: 'logs/20210510-184646'
```

```
In [ ]: #After VGG-16 network without FC layers, add a new Conv block ( 1 Conv layer and 1 Maxpooling ), 2 FC layers and a ou
# You are free to choose any hyperparameters/parameters of conv block, FC layers, output layer

#Input layer
input_layer = Input(shape=(224,224,3),name='Input_Layer')
vgg16_layer= vgg_(input_layer)

#vgg16_layer1 =vgg16_layer.layers[-1].output
#Conv Layer
Conv1 = Conv2D(filters=4096,kernel_size=(7,7),strides=(1,1),padding='valid',
               activation='relu',kernel_initializer=tf.keras.initializers.he_normal(seed=25),name='Conv1')(vgg16_layer

#Conv Layer
Conv2 = Conv2D(filters=4096,kernel_size=(1,1),strides=(1,1),padding='valid',
               activation='relu',kernel_initializer=tf.keras.initializers.he_normal(seed=27),name='Conv2')(Conv1)

#Conv Layer
Conv3 = Conv2D(filters=4096,kernel_size=(1,1),strides=(1,1),padding='valid',
               activation='relu',kernel_initializer=tf.keras.initializers.he_normal(seed=27),name='Conv3')(Conv2)

#Flatten
flatten = Flatten(name='Flatten')(Conv3)
```

```

drop=Dropout(0.5)(flatten)

#output layer
Out = Dense(units=16,activation='softmax',kernel_initializer=tf.keras.initializers.glorot_normal(seed=29),name='Output')

#Creating a model
model2 = Model(inputs=input_layer,outputs=Out)
#model1 = Model(inputs = vgg16_layer, outputs = Out)

```

In []: `model2.summary()`

Model: "model_1"

Layer (type)	Output Shape	Param #
Input_Layer (InputLayer)	[(None, 224, 224, 3)]	0
vgg16 (Functional)	(None, 7, 7, 512)	14714688
Conv1 (Conv2D)	(None, 1, 1, 4096)	102764544
Conv2 (Conv2D)	(None, 1, 1, 4096)	16781312
Conv3 (Conv2D)	(None, 1, 1, 4096)	16781312
Flatten (Flatten)	(None, 4096)	0
dropout_1 (Dropout)	(None, 4096)	0
Output (Dense)	(None, 16)	65552
Total params: 151,107,408		
Trainable params: 136,392,720		
Non-trainable params: 14,714,688		

Tensorboard 2

In []: `#compiling`
`model2.compile(optimizer=tf.keras.optimizers.Adam(lr=0.0001),loss='categorical_crossentropy',metrics=['accuracy'])`

```
In [ ]: STEP_SIZE_TRAIN=train_generator.n//train_generator.batch_size
STEP_SIZE_VALID=valid_generator.n//valid_generator.batch_size
```

```
model2.fit(train_generator,
           steps_per_epoch=STEP_SIZE_TRAIN,
           validation_data=valid_generator,
           validation_steps=STEP_SIZE_VALID,
           epochs=7,
           callbacks=[tensorboard_callback])
```

```
Epoch 1/7
672/672 [=====] - 371s 546ms/step - loss: 1.7567 - accuracy: 0.4742 - val_loss: 1.0914 - val
_accuracy: 0.6680
Epoch 2/7
672/672 [=====] - 366s 545ms/step - loss: 0.9776 - accuracy: 0.7012 - val_loss: 0.9414 - val
_accuracy: 0.7174
Epoch 3/7
672/672 [=====] - 368s 547ms/step - loss: 0.7708 - accuracy: 0.7641 - val_loss: 0.9326 - val
_accuracy: 0.7181
Epoch 4/7
672/672 [=====] - 366s 545ms/step - loss: 0.6365 - accuracy: 0.8045 - val_loss: 0.8730 - val
_accuracy: 0.7370
Epoch 5/7
672/672 [=====] - 367s 545ms/step - loss: 0.5306 - accuracy: 0.8360 - val_loss: 0.8704 - val
_accuracy: 0.7505
Epoch 6/7
672/672 [=====] - 366s 545ms/step - loss: 0.4436 - accuracy: 0.8607 - val_loss: 0.8555 - val
_accuracy: 0.7638
Epoch 7/7
672/672 [=====] - 366s 545ms/step - loss: 0.3660 - accuracy: 0.8859 - val_loss: 1.0525 - val
_accuracy: 0.7383
```

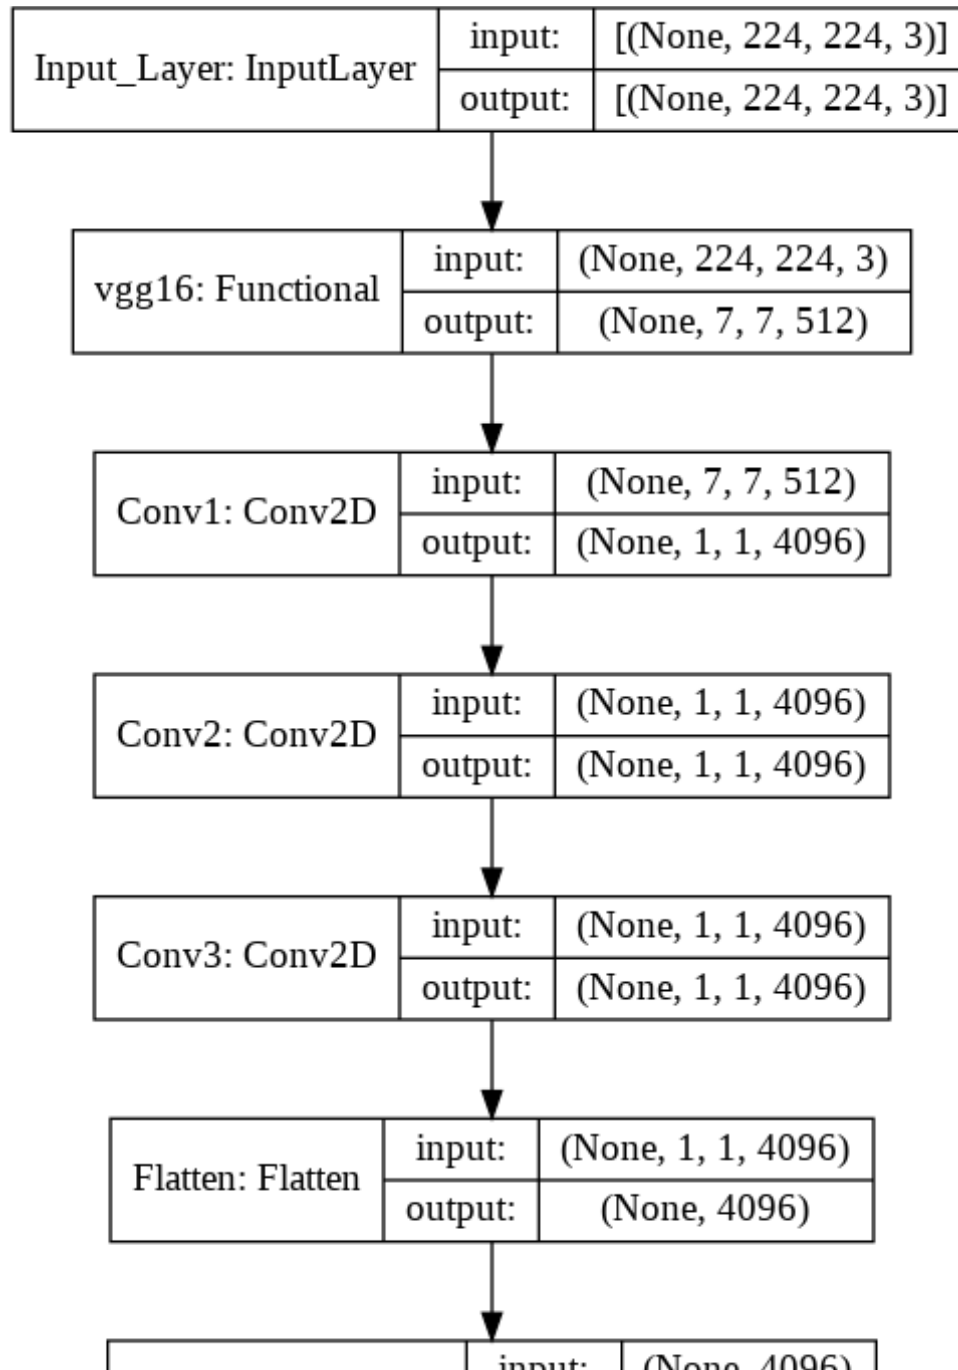
```
Out[ ]: <tensorflow.python.keras.callbacks.History at 0x7f097fa20810>
```

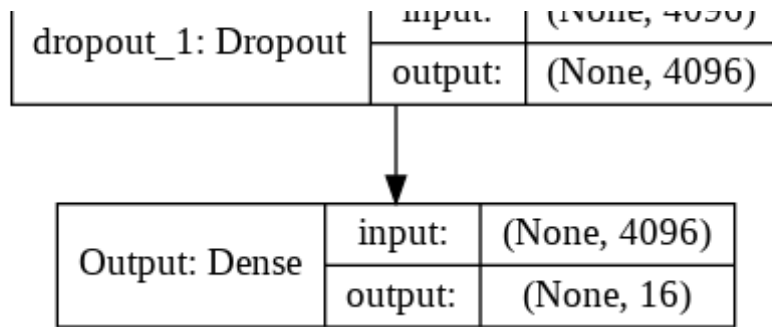
```
In [ ]: %load_ext tensorboard
%tensorboard --logdir $logdir
```

```
The tensorboard extension is already loaded. To reload it, use:
%reload_ext tensorboard
Reusing TensorBoard on port 6007 (pid 525), started 0:59:59 ago. (Use '!kill 525' to kill it.)
```

```
In [ ]: from tensorflow.keras.utils import plot_model
plot_model(model2, 'model2.png', show_shapes=True)
```

Out[]:





Model 3

```
In [ ]: from keras.models import Sequential
from tensorflow.keras.models import Model
from keras.layers import Dense, Activation, Flatten, Dropout, Input
from keras.layers import Conv2D, MaxPooling2D
from keras import regularizers, optimizers
from keras.applications.vgg16 import VGG16
from tensorflow.keras.callbacks import ReduceLRonPlateau, EarlyStopping, ModelCheckpoint
import datetime
```

```
In [ ]: tf.keras.backend.clear_session()
```

```
In [15]: import os
logdir3 = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
tensorboard_callback3 = tf.keras.callbacks.TensorBoard(logdir3, histogram_freq=1, write_graph=True, write_grads=True)

WARNING:tensorflow:`write_grads` will be ignored in TensorFlow 2.0 for the `TensorBoard` Callback.
```

```
In [13]: #Use VGG-16 pretrained network without Fully Connected layers and initialize all the weights with Imagenet trained weights
vgg_3=VGG16(weights='imagenet',include_top=False,input_shape=(224,224,3))

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5
58892288/58889256 [=====] - 0s 0us/step
```

```
In [14]: vgg_3.trainable
```

```
Out[14]: True
```

```
In [35]: for layer in vgg_3.layers[:-6]: # only last 6 layers will be trainable
        layer.trainable=False
```

```
In [50]: # for layer in vgg_3.layers:
        # print(layer.name , "=",layer.trainable)
```

```
In [45]: #After VGG-16 network without FC layers, add a new Conv block ( 1 Conv layer and 1 Maxpooling ), 2 FC layers and a ou
        # You are free to choose any hyperparameters/parameters of conv block, FC layers, output layer
```

```
#input layer
input=Input(shape=(224,224,3))
vgg_layers=vgg_3(input)

#conv layer
layer1=Conv2D(filters=4096 ,kernel_size=(7,7) ,strides=(1, 1),padding='valid',
              activation='relu',kernel_initializer=tf.keras.initializers.he_normal(seed=23))(vgg_layers)

#conv layer
layer2=Conv2D(filters=4096 ,kernel_size=(1,1) ,strides=(1, 1),padding='valid',
              activation='relu',kernel_initializer=tf.keras.initializers.he_normal(seed=28))(layer1)

#flatten
layer3=Flatten()(layer2)
#drop=Dropout(0.5)(flatten)

#output layer
output=Dense(16,activation='softmax',kernel_initializer=tf.keras.initializers.glorot_normal(seed=39))(layer3)

#creating a model
model3 = Model(inputs=input, outputs=output)
```

```
In [46]: model3.compile(optimizer=tf.keras.optimizers.Adam(lr=0.001),loss='categorical_crossentropy',metrics=['accuracy'])
```

```
In [47]: model3.summary()
```

Model: "model_12"

Layer (type)	Output Shape	Param #
input_6 (InputLayer)	[(None, 224, 224, 3)]	0

vgg16 (Functional)	(None, 7, 7, 512)	14714688
conv2d_11 (Conv2D)	(None, 1, 1, 4096)	102764544
conv2d_12 (Conv2D)	(None, 1, 1, 4096)	16781312
flatten_1 (Flatten)	(None, 4096)	0
dense_1 (Dense)	(None, 16)	65552
=====		
Total params: 134,326,096		
Trainable params: 129,050,640		
Non-trainable params: 5,275,456		

```
In [49]: #STEP_SIZE_TRAIN=train_generator.n//train_generator.batch_size
#STEP_SIZE_VALID=valid_generator.n//valid_generator.batch_size
```

```
model3.fit_generator(train_generator,
steps_per_epoch=33600/50, #STEP_SIZE_TRAIN
validation_data=valid_generator,
validation_steps=14400/50, #STEP_SIZE_VALID
epochs=4,
callbacks=[tensorboard_callback3])
```

```
/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/engine/training.py:1844: UserWarning: `Model.fit_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.
```

```
warnings.warn("`Model.fit_generator` is deprecated and ")
```

```
Epoch 1/4
```

```
672/672 [=====] - 505s 752ms/step - loss: 2.7728 - accuracy: 0.0629 - val_loss: 2.7729 - val_accuracy: 0.0593
```

```
Epoch 2/4
```

```
672/672 [=====] - 504s 750ms/step - loss: 2.7727 - accuracy: 0.0617 - val_loss: 2.7730 - val_accuracy: 0.0593
```

```
Epoch 3/4
```

```
672/672 [=====] - 506s 753ms/step - loss: 2.7727 - accuracy: 0.0643 - val_loss: 2.7731 - val_accuracy: 0.0593
```

```
Epoch 4/4
```

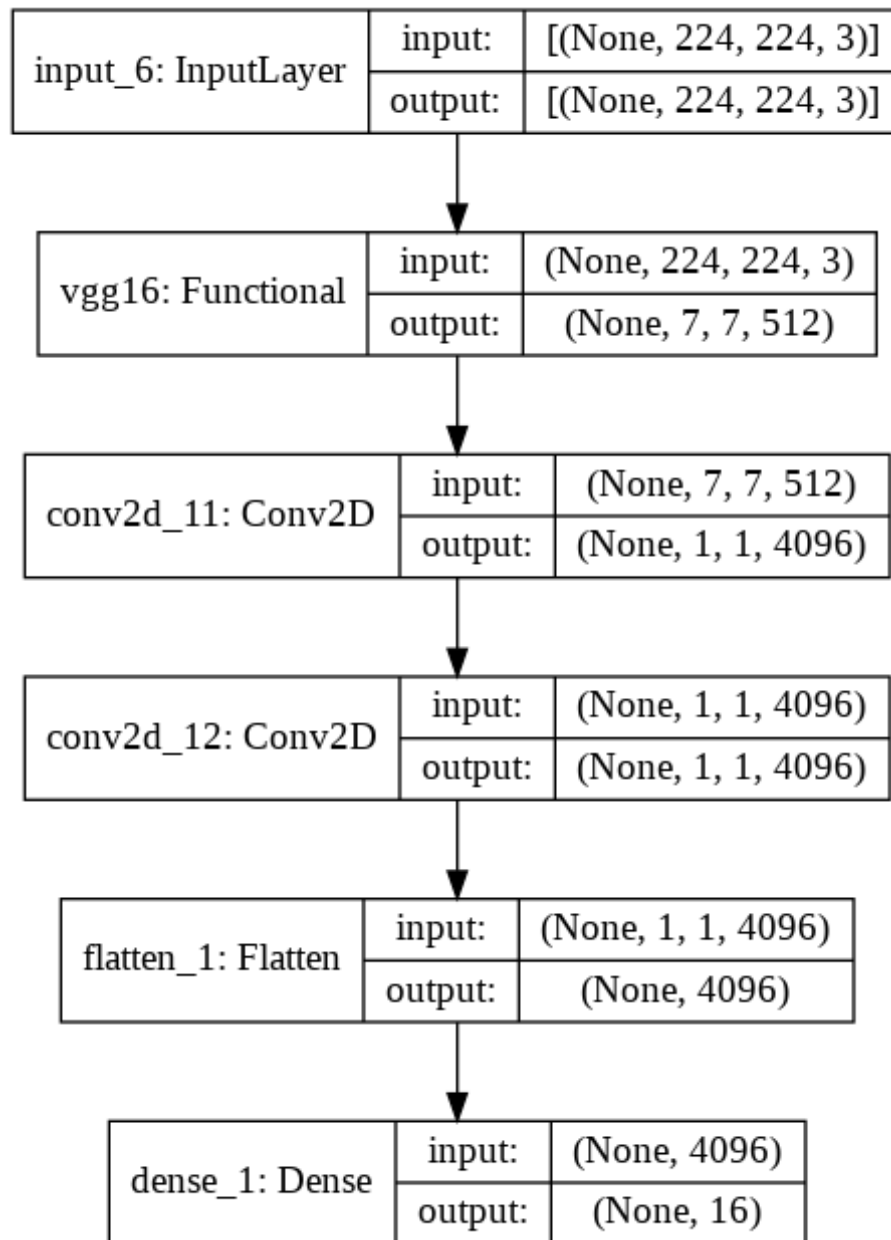
```
672/672 [=====] - 506s 753ms/step - loss: 2.7727 - accuracy: 0.0640 - val_loss: 2.7731 - val_accuracy: 0.0593
```

```
Out[49]: <tensorflow.python.keras.callbacks.History at 0x7f570a29fa50>
```

```
In [51]: %load_ext tensorboard
%tensorboard --logdir $logdir3
```

```
In [52]: from tensorflow.keras.utils import plot_model
plot_model(model3, 'model2.png', show_shapes=True)
```

Out[52]:



In []:

